



4. QUERY PROCESSING PART 1

Slides adapted from
Pearson Ed.

In this lecture – learning objectives

1

- Understand how a computer is able to translate natural-language SQL queries into a set of actions to perform on data.

2

- **Query decomposition:** Be familiar with both generic parsing techniques and ones specific to SQL queries / relational databases.

3

- **Query optimization:** Multiple ways exist to execute a single query.
 - Recognize equivalent relational algebra expressions.
 - Distinguish between slower and faster executions of the same query.

Query processing in RDBMSs

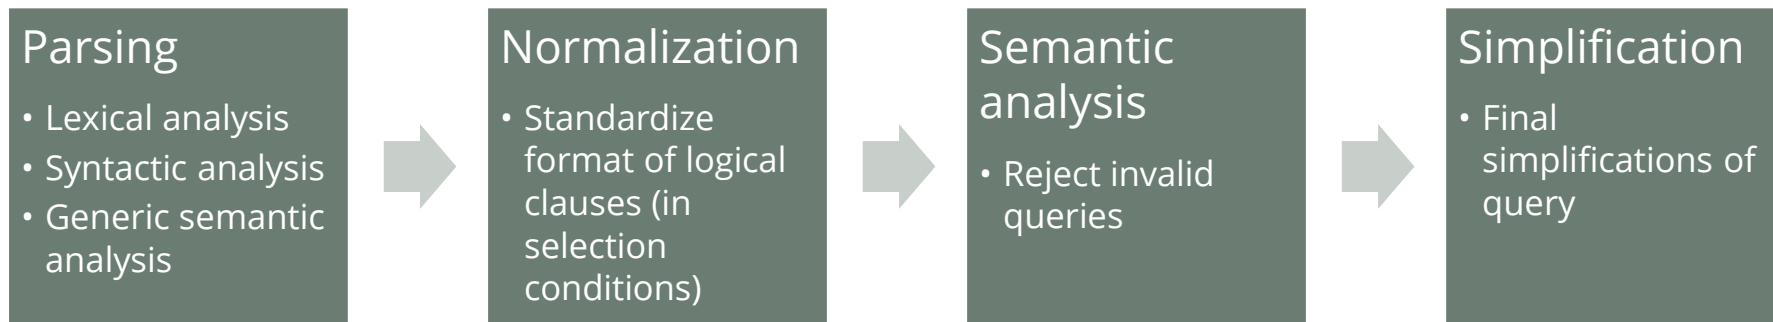
- SQL is a declarative (rather than procedural) language – specifies the desired information but not the strategy for retrieval.
- Job of optimizing performance of queries can be given to / centralized in a DBMS.
- Two components of query processing:
 - **Query decomposition:** conversion of SQL's natural language syntax into equivalent and valid relational algebra.
 - **Query optimization:** selection between equivalent relational algebra statements for the one that minimizes work (usually disk reads).



Query decomposition

Query decomposition

- Goal is to reject invalid SQL statements and convert valid SQL statements into equivalent relational algebra statement.
- Stages:
 - Parsing, normalization, semantic analysis, simplification.



Parsing

- SQL statement parsed using compiler techniques:
 - Lexical analysis to understand tokens (keywords, variable names, comma-separated lists, etc.).
 - Lexicon = dictionary. Make sure all parts of a query are a recognizable token.
 - Syntactic analysis to understand arrangement of tokens and check that they are grammatically valid.
 - Syntax = grammar / how words are properly arranged into sentences.
 - Generic semantic analysis to detect name mismatches.
 - Semantics = meaning. Is the query trying to do something sensible?

Parsing example

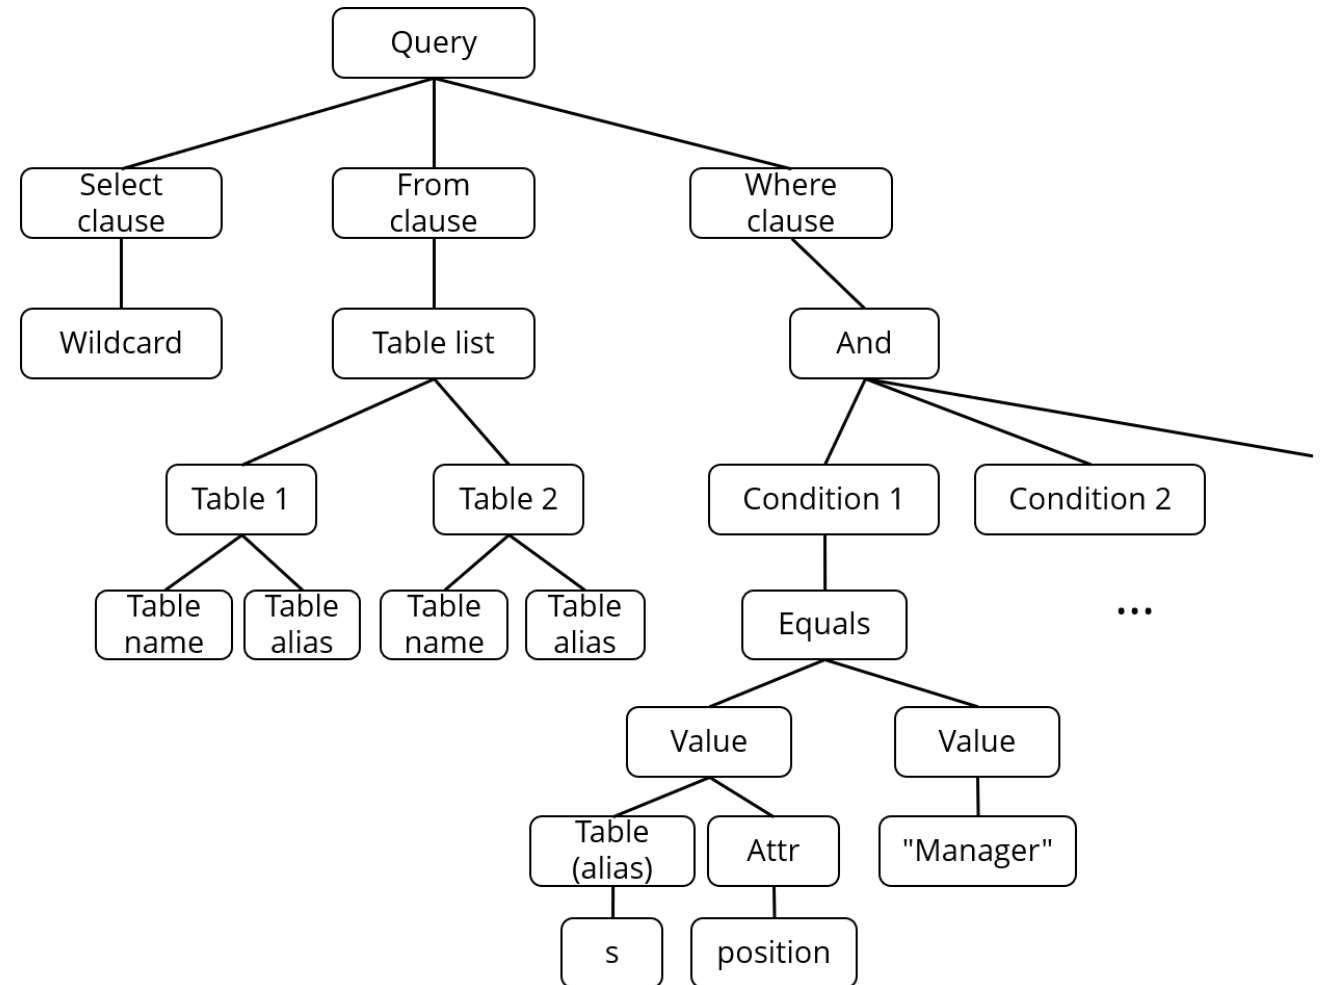
- Break query into tokens.
- Use a grammar to distinguish between acceptable and non-acceptable combinations of tokens.
- Flag issues and report them back to the user.

SELECT *
FROM Staff s, Branch b
WHERE s.position = "Manager"
AND b.city = "London"
AND s.branchNo = b.branchNo;

SELECT keyword WILDCARD
FROM keyword VARIABLE VARIABLE ALIAS COMMA VARIABLE VARIABLE ALIAS
WHERE keyword VARIABLE DOT ATTR EQUALS STRING LITERAL
AND keyword VARIABLE DOT ATTR EQUALS STRING LITERAL
AND keyword VARIABLE DOT ATTR EQUALS VARIABLE DOT ATTR SEMI-COLON

Result: abstract syntax tree

- Summarizes structure of query.

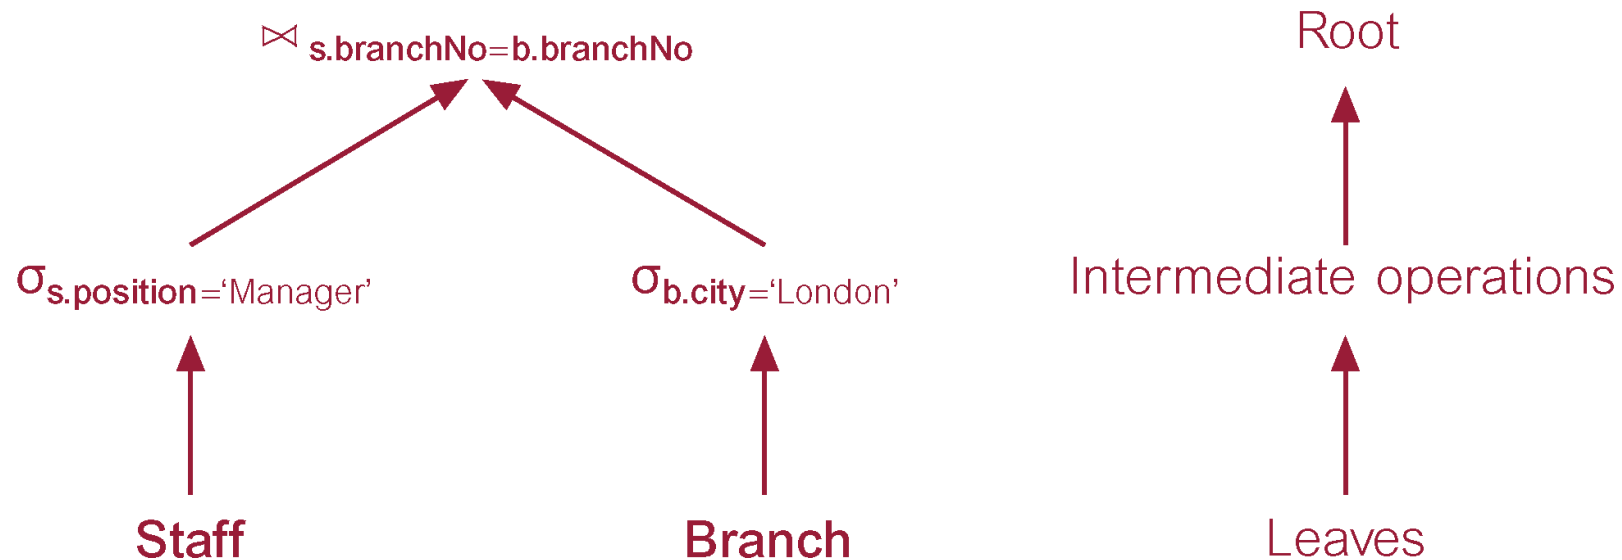


Parsing: checking for errors

- Generic semantic analysis:
 - Check that identifiers (table and attribute names) are valid.
 - Verify operations are appropriate to apply to types of objects.
- Example:
SELECT staffNumber FROM Staff WHERE position > 10;
 - Issues:
 - "staffNumber" not defined – should be "staffNo".
 - Numerical comparison inappropriate for string data of "position".
- If everything looks okay, can convert from syntax tree to a relational algebra tree.

Relational algebra tree

- Leaf = relation; interior node = RA op; root = final op of query; arrows = connect args to op.



Normalization

- Can convert conditions into a standard form that makes checking them easier.
 - Conjunctive normal form: all conditions are made up of “OR” clauses “AND”ed together (e.g. $(p \vee q) \wedge r \wedge (s \vee t \vee w)$).
 - Disjunctive normal form: all conditions are made up of “AND” clauses “OR”ed together (e.g. $(p \wedge q) \vee r \vee (s \wedge t \wedge w)$).
- Helpful for semantic analysis (next step).

Semantic analysis


- Rejects normalized conditional formulas that are incorrectly formulated or contradictory.
 - Incorrectly formulated: components do not contribute to generation of result.
 - Contradictory: formula cannot be satisfied by any tuple.

Semantic analysis: Checking for incorrect formulations

- Draw *relation connection graph* with one node per relation in query and one node for the result.
- Add edges between relations if they're participating in a join condition together.
- Add edges between result node and relations that appear in final list of attributes.
- Unconnected graph = incorrect formulation.

Relation connection graph example

- ```
SELECT p.propertyNo, p.street
FROM Client c, Viewing v,
 PropertyForRent p
WHERE c.clientNo = v.clientNo AND
 c.maxRent >= 500 AND
 c.prefType = 'Flat' AND
 p.ownerNo = 'CO93';
```

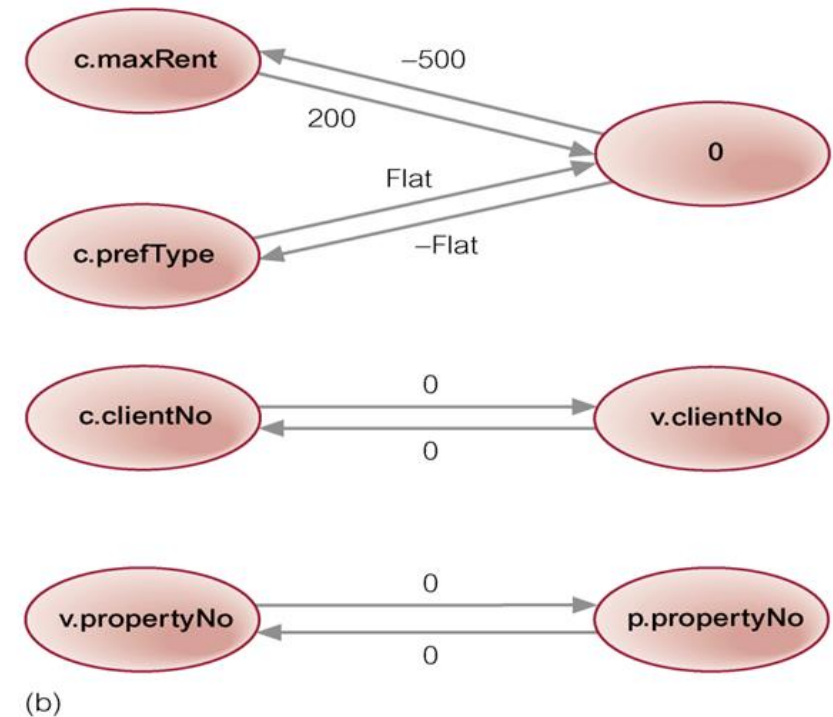


# Semantic analysis: Checking for contradictions

- Draw a *normalized attribute connection graph* with one node for each attribute and a '0' constant node.
- Add edges from attribute a to attribute b if they're involved in an ' $a \leq b$ ' comparison, and between an attribute and the '0' node if the attribute is compared to a constant.
  - Equality is represented with two arrows.
  - Use weights on arrows to adjust the constant appropriately.
- Cycle with negative total weight = unsatisfiable condition.

# Normalized attribute connection graph

- SELECT p.propertyNo, p.street  
FROM Client c, Viewing v,  
PropertyForRent p  
WHERE c.clientNo = v.clientNo AND  
c.maxRent >= 500 AND  
v.propertyNo = p.propertyNo AND  
c.prefType = 'Flat' AND  
c.maxRent < 200;





# Simplification

- Query seems okay.
- Check if user has permissions to access everything it touches.
  - Reject operation otherwise.
- Simplify conditions using Boolean algebra laws and knowledge of integrity constraints.



# Query optimization

# Example of importance of query optimization

- Find all Managers who work at a London branch:  
`SELECT *`  
`FROM Staff s, Branch b`  
`WHERE s.branchNo = b.branchNo AND (s.position = 'Manager'`  
`AND b.city = 'London');`

# Different strategies

■ Three equivalent RA queries are:

1.  $\sigma_{(\text{position}='Manager') \wedge (\text{city}='London') \wedge (\text{Staff.branchNo}=\text{Branch.branchNo})} (\text{Staff} \times \text{Branch})$
2.  $\sigma_{(\text{position}='Manager') \wedge (\text{city}='London')} (\text{Staff} \bowtie_{\text{Staff.branchNo}=\text{Branch.branchNo}} \text{Branch})$
3.  $(\sigma_{\text{position}='Manager'}(\text{Staff})) \bowtie_{\text{Staff.branchNo}=\text{Branch.branchNo}} (\sigma_{\text{city}='London'}(\text{Branch}))$

# Cost comparison

- Assuming:
  - 1000 staff members, 50 branches, 50 managers, 5 London branches.
  - No search aids such as indices.
  - Tuples accessed one at a time.
- Cost in disk accesses (reads/writes):
  1. Read Staff and Branch, write full Cartesian product, comb through result to select = 101,050 disk accesses.
  2. Read Staff and Branch, write results of joining Staff to matching Branch info, comb through result set = 3,050 disk accesses.
  3. Read/select Staff, write manager set, read/select Branch, write London set, read manager/London sets to join = 1,160 disk accesses.

# Requirements for query optimization

- An understanding of the space of equivalent relational algebra statements to create alternatives.
- Strategy for selecting more efficient options.
- Aided by **statistics** to more accurately approximate the cost of operations.

# Two strategies for query optimization: Dynamic vs. static processing

- Possible to process query each time the query is run (*dynamic query optimization*).
  - Uses most up-to-date information/statistics.
  - Time-consuming and less responsive to the user.
- Alternatively, possible to pre-process query and use pre-decided strategy when query is run (*static query optimization*).
  - Doesn't take into account current statistics.
  - More optimization possible because response time isn't an issue.
- Hybrid approaches: re-optimization triggered by significant database changes, or compute once per session and cache.



# Relational algebra transformation rules



# Equivalences

- Intersection:  $R \cap S = R - (R - S)$
- Join:  $R \bowtie S = \sigma_{predicate}(R \times S)$
- Natural join:  $\Pi_{reduced\ list}(\sigma_{R.c1=S.c1,etc.}(R \times S))$
- Etc.

# Conjunctions in selection conditionals

- Conjunctions in a selection condition can be broken into one selection per clause:
  - $\sigma_{p \wedge q}(R) = \sigma_p(\sigma_q(R))$
- And reordered (commutativity):
  - $\sigma_p(\sigma_q(R)) = \sigma_q(\sigma_p(R))$

# Projection simplification

- A series of valid projections can be reduced to just the last (outer) projection.
  - Valid = assumes that it does not try to project onto missing / non-existent attributes.
  - $\Pi_{L_1} \Pi_{L_2} \dots \Pi_{L_n}(R) = \Pi_{L_1}(R)$

# Distributing selection across a projection

- Assuming a selection's condition involves a subset of a projection's attributes, can freely flip the order of the two operations:
  - $\Pi_A(\sigma_p(R)) = \sigma_p(\Pi_A(R))$  where  $\text{attr}(p) \subseteq A$
- Cannot do this if selection involves attributes not in the projection, as performing the selection after the projection is invalid.

# Commutation properties of binary operations

- The following operations commute (are symmetric):
  - Cartesian product:  $R \times S = S \times R$
  - Joins (including equijoin, natural join):  $R \bowtie_p S = S \bowtie_p R$
  - Set union:  $R \cup S = S \cup R$
  - Set intersection:  $R \cap S = S \cap R$
  - But not set difference.
- Switching order of products/joins changes order of attributes, but does not change relation overall.

# Distributing selection across binary operations

- Can distribute a selection across:
  - Any set operation including set difference:
$$\sigma_p(R \cup S) = \sigma_p(R) \cup \sigma_p(S)$$
$$\sigma_p(R \cap S) = \sigma_p(R) \cap \sigma_p(S)$$
$$\sigma_p(R - S) = \sigma_p(R) - \sigma_p(S)$$
  - Cartesian products and joins as long as selection condition clauses can be separated into ones that only apply to one relation or the other
$$\sigma_{p \wedge q}(R \times S) = \sigma_p(R) \times \sigma_q(S)$$
$$\sigma_{p \wedge q}(R \bowtie S) = \sigma_p(R) \bowtie \sigma_q(S)$$
(p only involves R's attrs; q only involves S's attrs)

# Distributing projection across binary operations

- Can distribute a projection across:

- Any set operation including set difference:

$$\Pi_A(R \cup S) = \Pi_A(R) \cup \Pi_A(S)$$

$$\Pi_A(R \cap S) = \Pi_A(R) \cap \Pi_A(S)$$

$$\Pi_A(R - S) = \Pi_A(R) - \Pi_A(S)$$

- Cartesian products:  $\Pi_A(R \times S) = \Pi_A(R) \times \Pi_A(S)$

- Joins as long as projection doesn't mask out attributes needed for the join condition clauses (if any):  $\Pi_{L_1 \cup L_2}(R_1 \bowtie_p R_2) = \Pi_{L_1 \cup L_2}(\Pi_{L_1 \cup A}(R_1) \bowtie_p \Pi_{L_2 \cup A}(R_2))$   
(L1= attrs from R1 in final projection; L2 = attrs from R2 in final projection; A = join condition attrs)

- If no conditional, can simplify to:  $\Pi_{L_1 \cup L_2}(R_1 \bowtie_p R_2) = \Pi_{L_1}(R_1) \bowtie_p \Pi_{L_2}(R_2)$

# Associative properties of binary operations

- The following operations are associative:
  - Set union and intersection (but not difference):
$$(R \cup S) \cup T = R \cup (S \cup T)$$
$$(R \cap S) \cap T = R \cap (S \cap T)$$
$$(R - S) - T \neq R - (S - T)$$
  - Cartesian product:  $(R \times S) \times T = R \times (S \times T)$
  - Natural joins:  $(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$
  - Other joins as long as join conditions are handled correctly to reflect that different pairs of relations are being joined
$$(R \bowtie_p S) \bowtie_q T = R \bowtie_{p'}, (S \bowtie_{q'} T)$$





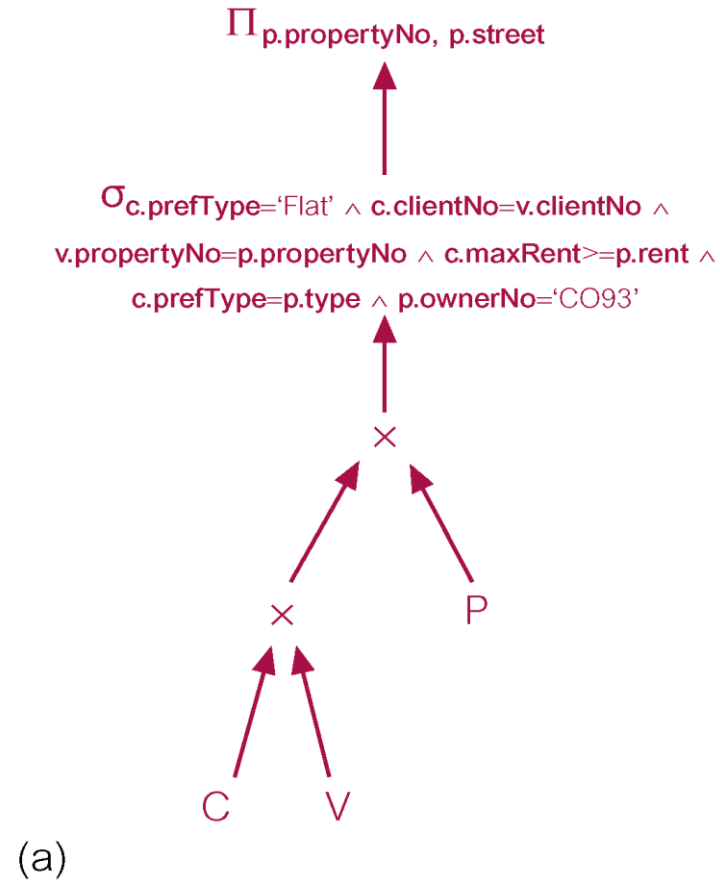
# Using RA rules for query optimization

# Example: Using transformation rules

- Task: For prospective renters of flats, find properties that match requirements and owned by owner CO93.


```
SELECT p.propertyNo, p.street
FROM Client c, Viewing v, PropertyForRent p
WHERE c.prefType='Flat' AND
 c.clientNo = v.clientNo AND
 v.propertyNo = p.propertyNo AND
 c.maxRent >= p.rent AND
 c.prefType = p.type AND
 p.ownerNo = 'CO93';
```

# Original RA tree from SQL query




# Distributing selections inward

- Can separate conjunctions of a selection.
- Can distribute selections across Cartesian products as long as they only use an individual relation's attrs in their condition.



# Join equivalences


- Can convert product + selection into equivalent equijoin



(c)


# Associativity of joins

- Can use associativity rules to reorder which join is performed first (innermost)
- Choose the most selective to go first (the one that selects all flats owned by one owner)



# Distributing projections

- Move projections inward.



# Heuristics for speeding up queries

- Perform selection and projection ops as early as possible because they reduce the size of datasets.
  - Distribute ops inward using rules.
- Combine Cartesian products followed by selections with joins to combine two ops into one.
- Perform most restrictive selections first.
  - Rearrange leaves using associativity rules for binary operations.
- Compute common expressions once and store result (if not too big).



# Summary

- Parsing techniques help a computer translate natural-language SQL queries into *relational algebra trees* (RATs) corresponding to an RA expression.
- Different RATs can have drastically different execution time.
- Use RA equivalences to transform an RAT into one that is more efficient but *does the same work*.